

Qualm AI

Product & Technical Documentation Report

Created: July 2026

Engineered for Strategic Clarity and Absolute Privacy

Qualm AI - Product & Technical Documentation Report

This document serves as the official, professional documentation for Qualm AI, detailing its core philosophy, product features, technical stack, security architecture, and design system.

1. Introduction to Qualm AI

The Philosophy: "Quiet Intelligence"

Most modern conversational AI platforms focus on velocity - generating words as fast as possible. Qualm AI is built on the opposite philosophy: deliberate reflection.

Qualm is a premium conversational partner designed for strategic logic, meticulous research, and contextual clarity. Instead of rushing to spit out answers, Qualm is engineered to evaluate parameters, question assumptions, critique its own reasoning steps, and verify information before displaying a response. It is the AI that **thinks before you act**.

Target Audience

Qualm AI is built for:

- **Decision Makers & Executives**: Who need to stress-test business expansion plans or organizational strategy.
- **Researchers & Analysts**: Who require citation-backed document synthesis rather than general summaries.
- **Privacy-Conscious Professionals**: Who want direct control over their data, API keys, and conversational history without cloud-broker databases.

2. Product Capabilities & Features

Qualm AI is split into two primary components: a Premium Introduction Landing Page and an Intelligent Discourse Web Client. Together, they offer the following core features:

Feature	Description	User Benefit
Reasoning Loops	Loops query parameters through validation steps, critiquing assumptions before generating answers.	Minimizes hallucinations and outputs highly structured, logical guides.
Memoir Synthesis	Automatically saves key corporate values, goals, and strategies into a continuous local memory.	Retains context across development cycles without losing session memory.
Decision Simulations	Runs commercial concept tests against simulated regional, logistical, and capacity stress points.	Isolates potential strategy failures before committing capital.
Meticulous Synthesis	Parses attached PDFs, financial spreadsheets, or txt	

logs with high citation accuracy.

Extracts data from complex documents into

clean markdown grids.

3. Technical Architecture & Tech Stack

Qualm AI is engineered as a serverless, client-first web application. This minimizes backend latency, eliminates database maintenance costs, and maximizes user data security.

The Technology Stack

1. Core Platform: HTML5, CSS3, and modern Vanilla JavaScript (ES6+).
2. Styling & Layout:
 - **TailwindCSS**: Utilized via CDN for rapid container and utility grids in the chat client UI.
 - **Custom Premium CSS**: Applied to build the landing page's custom glassmorphism components, overlay blobs, and particle canvas rendering.
3. AI Models Integration:
 - **Gemini 2.5 Flash**: Default model for lightweight, fast strategic tasks.
 - **Gemini 2.5 Pro**: Analytical model for complex coding, file synthesis, and logic critiques.
4. Authentication: Integrated with Google Identity Services (OAuth2) for secure client authentication.
5. Android Wrapper: Wrapped via AppsGeyser to compile the web application into an installable Android package (.apk).

4. Privacy & Security Infrastructure

Data security is the cornerstone of Qualm AI. Unlike other AI agents that save prompts on central databases to train future models, Qualm relies on a zero-trust, local security model.

- **Direct Endpoint Routing**: Prompts are sent directly from your browser to Google's Gemini API servers. No intermediary servers intercept or log your conversations.
- **Zero Training Data Policy**: Using your custom Gemini API keys ensures your queries are treated as private enterprise transactions. Your corporate secrets are never used to train public LLM models.
- **Client-Side Cryptography & Storage**:
 - API keys, settings, user avatars, and discourse logs are saved directly in the browser's `localStorage``.
 - Conversations and credentials never touch a Qualm cloud server. Users retain absolute ownership and can clear the cache in one click.

5. UI/UX Design System

Qualm's visual design is crafted to look premium, minimal, and state-of-the-art.

Spacing & Layout

The page follows a structured grid layout with optimized vertical spaces (60px margins/paddings on sections) to balance whitespace and content flow.

Color Palette

The aesthetic utilizes a curated, warm color palette that avoids generic digital styling:

- **Cream** (`#FDFCF0`): Used as the primary background color. Sleek, warm, and easy on the eyes.
- **Charcoal** (`#1A1A1A`): Used for bold typography, primary buttons, and system headers.
- **Stone** (`#4A4A4A`): Used for secondary paragraph text and labels.
- **Glow Accents**: Subtle pastel blobs (`#e3b5ff` purple, `#adc7ff` blue, `#00dbe7` cyan) float in the background to add depth.

Typography

Two distinct typography styles are paired to create an editorial feel:

- **Instrument Serif** (Google Fonts): Used for logo labels, editorial headings, and quotes to convey premium elegance.
- **Hanken Grotesk / Inter** (Google Fonts): A clean, modern sans-serif typeface used for paragraphs and technical information to ensure maximum readability.

6. Installation & Client Distribution

Qualm AI is accessible across multiple client platforms:

1. Web Client: Access the active conversational workspace immediately via qu-tqib.onrender.com.
2. Android Application:
 - Compiled into a mobile package hosted at AppsGeyser: [\[https://appsgeyser.io/19958856/Qualm\]](https://appsgeyser.io/19958856/Qualm)(<https://appsgeyser.io/19958856/Qualm>).
 - Installable via QR Code scan directly from the landing page's download center.
3. Desktop Apps: Installable across Windows (setup.exe), macOS (.dmg), and Linux (.deb/.rpm) through simulated wrappers.